

Text Preprocessing

COMP90042

Natural Language Processing

Lecture 2



THE UNIVERSITY OF
MELBOURNE

Definitions

- Corpus: a collection of documents.
- Document: one or more sentences.
- Sentence
 - ▶ “The student is enrolled at the University of Melbourne.”
- Words
 - ▶ Sequence of characters with a meaning and/or function
- Word **token**: each instance of “the” in the sentence above.
 - E.g. 9 word tokens in the example sentence.
- Word **type**: the distinct word “the”.
 - ▶ Lexicon (“dictionary”): a group of word types.
 - ▶ E.g. 8 word types in the example sentence.

How Many Unique Words?

	#Tokens (N)	#Type (IVI)
Switchboard phone conversation	2.4 million	20 thousand
Shakespeare	800 thousand	31 thousand
Google N-gram	1 trillion	13 million

Church and Gale (1990): $IVI > O(N^{1/2})$

Why Preprocess?

- Most NLP applications have documents as inputs:
 - ▶ “This movie is so great!!! U should definitely watch it in the theater! Best sci-fi eva!” → 😊
 - ▶ “Eu estive em Melbourne no ano passado.” → “I was in Melbourne last year.”
- **Key point:** language is **compositional**. As humans, we can break these documents into individual components. To understand language, a computer should do the same.
- **Preprocessing** is the first step.

Preprocessing Steps

1. Remove unwanted formatting (e.g. HTML)
2. **Sentence segmentation**: break documents into sentences
3. **Word tokenisation**: break sentences into words
4. **Word normalisation**: transform words into canonical forms
5. **Stopword removal**: delete unwanted words

“<p> Hi there. I’m
TARS. </p>”

["Hi there.",
"I'm TARS."]

[["hi", "there", "."],
["i", "am", "tars", "."]]

“Hi there. I’m
TARS.”

[["Hi", "there", "."],
["I", "m", "TARS", "."]]

[[], ["tars"]]

Sentence Segmentation

Sentence Segmentation

- Naïve approach: break on sentence punctuation ([.?!])
 - ▶ But periods are used for abbreviations!
(U.S. dollar, ..., Yahoo! as a word)
- Second try: use regex to require capital ([.?!] [A-Z])
 - ▶ But abbreviations often followed by names (Mr. Brown)
- Better yet: have lexicons
 - ▶ But difficult to enumerate all names and abbreviations
- State-of-the-art uses machine learning, not rules

Binary Classifier

- Looks at every “.” and decides whether it is the end of a sentence.
 - ▶ Decision trees, logistic regression
- Features
 - ▶ Look at the words before and after “.”
 - ▶ Word shapes:
 - Uppercase, lowercase, ALL_CAPS, number
 - Character length
 - ▶ Part-of-speech tags:
 - Determiners tend to start a sentence

Word Tokenisation

Word Tokenisation: English

- Naïve approach: separate out alphabetic strings (`\w+`)
- Abbreviations (*U.S.A.*)
- Hyphens (*merry-go-round* vs. *well-respected* vs. *yes-but*)
- Numbers (*1,000,00.01*)
- Dates (*3/1/2016*)
- Clitics (*n't* in *can't*)
- Internet language (*http://www.google.com*, *#metoo*, *:-*)
- Multiword units (*New Zealand*)

Word Tokenisation: Chinese

- Some Asian languages are written without spaces between words
- In Chinese, words often correspond to more than one character

墨大 的 学生 与众不同

Unimelb 's students (are) special

Word Tokenisation: Chinese

- Standard approach assumes an existing vocabulary
- MaxMatch algorithm
 - ▶ Greedily match longest word in the vocabulary

$V = \{\text{墨, 大, 的, 学, 生, 与, 众, 不, 同, 墨大, 学生, 不同, 与众不同}\}$

墨大的学生与众不同

match 墨大, match 的, match 学生, match 与众不同,
move to 的 move to 学 move to 与 done

Word Tokenisation: Chinese

- But how do we know what the vocabulary is
- And doesn't always work

not working

去 买

go buy

新西兰

New Zealand

花

flowers

去 买

go buy


新

new

西兰花

broccoli

Word Tokenisation: German

- *Lebensversicherungsgesellschaftsangestellter*
- = *life insurance company employee*
- Requires  *compound splitter*

Subword Tokenisation

- *Colourless green ideas sleep furiously* →
[colour] [less] [green] [idea] [s] [sleep] [furious] [ly]
- One popular algorithm: byte-pair encoding (BPE)
- Core idea: iteratively merge frequent pairs of characters
- Advantage:
 - ▶ Data-informed tokenisation
 - ▶ Works for different languages
 - ▶ Deals better with unknown words

Byte-Pair Encoding

- Dictionary

- ▶ [5] l o w _
- ▶ [2] l o w e s t _
- ▶ [6] n e w e r _
- ▶ [3] w i d e r _
- ▶ [2] n e w _

BPE.

- Vocabulary

- ▶ _, d, e, i, l, n, o, r, s, t, w

Byte-Pair Encoding

- Dictionary

- ▶ [5] l o w _

- ▶ [2] l o w e s t _

- ▶ [6] n e w e

- ▶ [3] w i d e

- ▶ [2] n e w _

end of word.

*9 times
most frequent*

put in vocab.

- Vocabulary

- ▶ _, d, e, i, l, n, o, r, s, t, w, r_

Byte-Pair Encoding

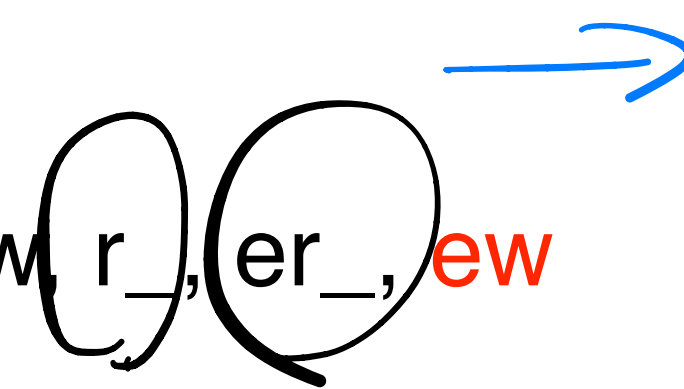
- Dictionary
 - ▶ [5] l o w _
 - ▶ [2] l o w e s t _
 - ▶ [6] n e w er_
 - ▶ [3] w i d er_
 - ▶ [2] n e w _
- Vocabulary
 - ▶ _, d, e, i, l, n, o, r, s, t, w, r_, er_

Byte-Pair Encoding

- Dictionary

- ▶ [5] l o w _
- ▶ [2] l o w e s t _
- ▶ [6] n ew er _
- ▶ [3] w i d er _
- ▶ [2] n ew _

- Vocabulary

- ▶ _, d, e, i, l, n, o, r, s, t, w, r_, er_, ew
- 

Byte-Pair Encoding

- Vocabulary

- ▶ `_`, `d`, `e`, `i`, `l`, `n`, `o`, `r`, `s`, `t`, `w`, `r_`, `er_`, `ew`
- ▶ `_`, `d`, `e`, `i`, `l`, `n`, `o`, `r`, `s`, `t`, `w`, `r_`, `er_`, `ew`, `new`
- ▶ `_`, `d`, `e`, `i`, `l`, `n`, `o`, `r`, `s`, `t`, `w`, `r_`, `er_`, `ew`, `new`, `lo`
- ▶ `_`, `d`, `e`, `i`, `l`, `n`, `o`, `r`, `s`, `t`, `w`, `r_`, `er_`, `ew`, `new`, `lo`, `low`
- ▶ `_`, `d`, `e`, `i`, `l`, `n`, `o`, `r`, `s`, `t`, `w`, `r_`, `er_`, `ew`, `new`, `lo`, `low`, `newer_`
- ▶ `_`, `d`, `e`, `i`, `l`, `n`, `o`, `r`, `s`, `t`, `w`, `r_`, `er_`, `ew`, `new`, `lo`, `low`, `newer_`, `low_`

*pick up
valid
words.*

Byte-Pair Encoding

- In practice BPE will run with thousands of merges, creating a large vocabulary
- Most frequent words will be represented as full words
- Rarer words will be broken into subwords
- In the worst case, unknown words in test data will be broken into individual letter

No unknown words.

Word Normalisation

Word Normalisation

- Lower casing (Australia → australia)
- Removing morphology ology
- Correcting spelling
- Expanding abbreviations (U.S.A → USA)
- Goal:
 - ▶ Reduce vocabulary
 - ▶ Maps words into the same type

Inflectional Morphology

- Inflectional morphology creates grammatical variants
- English inflects nouns, verbs, and adjectives
 - ▶ Nouns: *number* of the noun (-s)
 - ▶ Verbs: *number* of the subject (-s), the *aspect* (-ing) of the action and the *tense* (-ed) of the action
 - ▶ Adjectives: *comparatives* (-er) and *superlatives* (-est)
- Many languages have much richer inflectional morphology than English
 - ▶ E.g. French inflects nouns for gender (*un chat, une chatte*)
not changing the word classes

Lemmatisation

- Lemmatisation means removing any inflection to reach the uninflected form, the *lemma*
 - ▶ speaking → speak
- In English, there are irregularities that prevent a trivial solution:
 - ▶ poked → poke (not pok)
 - ▶ stopping → stop (not stopp)
 - ▶ watches → watch (not watche)
 - ▶ was → be (not wa)
- A lexicon of lemmas needed for accurate lemmatisation

Derivational Morphology

- Derivational morphology creates distinct words
- English derivational *suffixes* often change the lexical category, e.g.
 - ▶ -ly (personal → personally) *fun* noun *to*
 - ▶ -ise (final → finalise) *verb*
 - ▶ -er (write → writer)
- English derivational *prefixes* often change the meaning without changing the lexical category
 - ▶ write → rewrite
 - ▶ healthy → unhealthy

Stemming more aggressive than

Stemming

lemmatisation

do { Inflect
De v.
both

Stemming strips off all suffixes, leaving a stem

- ▶ E.g. automate, automatic, automation → automat
- ▶ Often not an actual lexical item
- Even less lexical sparsity than lemmatisation
- Popular in information retrieval
- Stem not always interpretable

(no ~~corpus~~)
lexicons
of lemmas

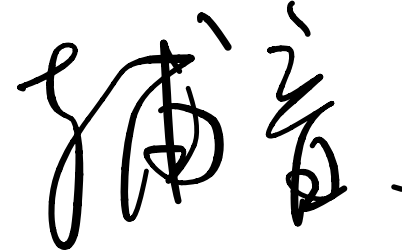
The Porter Stemmer

- Most popular stemmer for English
- Applies rewrite rules in stages
 - ▶ First strip inflectional suffixes
 - E.g. *-ies* → *-i*
 - ▶ Then derivational suffixes
 - E.g. *-isation* → *-ise* → *-i*

A sort of
stemming.

Rule based

The Porter Stemmer

- c (lowercase) = consonant; e.g. 'b', 'c', 'd'
- v (lowercase) = vowel; e.g. 'a', 'e', 'i', 'o', 'u'
- C = a sequence of consonants 
- ▶ s, ss, tr, bl
- V = a sequence of vowels
- ▶ o, oo, ee, io

The Porter Stemmer

- A word has one of the four forms:

▶ CVCV ... C

▶ CVCV ... V

▶ VCVC ... C

▶ VCVC ... V

optional [].

- Which can be represented as:

▶ [C]VCVC ... [V]

▶ [C] (VC)^m [V]

▶ m = measure

no repeating VC in the middle.

The Porter Stemmer



- $m=0$: TR, EE, TREE, Y, BY
 - $m=1$: TROUBLE, OATS, TREES, IVY
 - $m=2$: TROUBLES, PRIVATE, OATEN, ORRERY
-
- $\text{TR}\text{EE} = C(\text{VC})^0V$
 - $\text{TR}\text{EE}\text{S} = C(\text{VC})^1$
 - $\text{TROU}\text{BL}\text{E}\text{S} = C(\text{VC})^2$

The Porter Stemmer

- Rules format: (condition) S1 → S2

- e.g. (m > 1) EMENT → null

- ▶ REPLACEMENT → REPLAC

- Always use the longest matching S1

- ▶ CARESSES → CARESS

- ▶ CARESS → CARESS

- ▶ CARES → CARE

Rules:

SSES → SS

IES → I

SS → SS

S → null

The Porter Stemmer

- Step 1: plurals and past participles

	Rule	Positive Example	Negative Example
a	SSES → SS	caresses → caress	
	IES → I	ponies → poni	
	SS → SS	caress → caress	
	S → null	cats → cat	
b	(m>0) EED → EE	agreed → agree	feed → feed m=0
	(*v*) ED → null *v* = stem has vowel	plastered → plaster	bled → bled
	(*v*) ING →	motoring → motor	sing → sing
b+	AT → ATE	conflat(ed) → conflate	
c	(*v*) Y → I	happy → happi	

The Porter Stemmer

- Step 2, 3, 4: derivational inflections

	Rule	Positive Example
2	(m>0) ATIONAL → ATE	relational → relate
	(m>0) TIONAL → TION	conditional → condition
	(m>0) ENCI → ENCE	valenci → valence
	(m>0) ANCI → ANCE	hesitanci → hesitance
3	(m>0) ICATE → IC	triplicate → triplic
	(m>0) ATIVE → null	formative → form
	(m>0) ALIZE → AL	formalize → formal
4	(m>1) AL → null	revival → reviv
	(m>1) ER → null	airliner → airlin
	(m>1) ATE → null	activate → activ

The Porter Stemmer

- Step 5: tidying up

	Rule	Positive Example
5	$(m > 1) E \rightarrow \text{null}$	probate \rightarrow probat
	$(m = 1 \text{ and not } *o) E \rightarrow \text{null}$ $*o = \text{stem ends cvc, and second c is not w, x or y (e.g. -WIL, -HOP)}$	cease \rightarrow ceas
	$(m > 1 \text{ and } *d \text{ and } *L)$ $\text{null} \rightarrow \text{single letter}$ $*d = \text{stem ends with double consonant (e.g. -TT)}$ $*L = \text{stem ends with 'l'}$	controll \rightarrow control

The Porter Stemmer

- computational → comput
 - *Stem* step 2: ~~AT~~IONAL → ATE: computate
 - step 4: ~~ATE~~ → null: comput
- computer → comput
 - step 4: ER → null: comput

Fixing Spelling Errors

- typos in vocab.

- Why fix them?

- ▶ Spelling errors create new, rare types
- ▶ Disrupt various kinds of linguistic analysis
- ▶ Very common in internet corpora
- ▶ In web search, particularly important in queries

- How?

- ▶ String distance (Levenshtein, etc.)
- ▶ Modelling of error types (phonetic, typing etc.)
- ▶ Use an n-gram language model

characters
level

→ trained out

Other Word Normalisation

- Normalising spelling variations
 - ▶ Normalize → Normalise (or vice versa)
 - ▶ U r so coool! → *you are so cool*
- Expanding abbreviations
 - ▶ US, U.S. → United States
 - ▶ imho → in my humble opinion

Stopword Removal

Stop Words

- Definition: a list of words to be removed from the document
 - ▶ Typical in bag-of-word (BOW) representations
 - ▶ Not appropriate when sequence is important
- How to choose them?
 - ▶ All closed-class or function words
 - E.g. *the, a, of, for, he, ...*
 - ▶ Any high frequency words
 - ▶ NLTK, spaCy NLP toolkits

A Final Word

- Preprocessing unavoidable in text analysis
- Can have a major effect on downstream applications
- Exact steps may vary depending on corpus, task
- Simple rule-based systems work well, but rarely perfectly
- Language-dependent

Further Reading

- J&M3 Ch 2. on Normalisation
(includes a review of regex and Levenshtien distance)
- Details on the Porter Stemmer algorithm <http://snowball.tartarus.org/algorithms/porter/stemmer.html>