

Formal Language Theory & Finite State Automata

COMP90042

Natural Language Processing
Lecture 13



What is a Language?

- Methods to process sequence of symbols:
 - ▶ Language Model
 - ▶ Hidden Markov Model
 - ▶ Recurrent Neural Networks
- Nothing is fundamentally linguistic about these models

Formal Language Theory

- A language = set of **strings**
- A string = sequence of **elements** from a finite **alphabet**

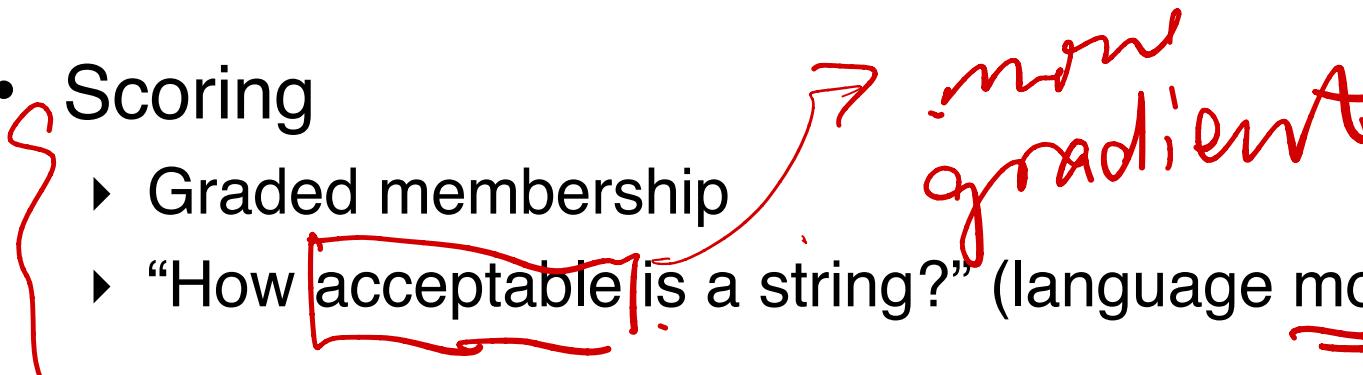
Motivation

- Formal language theory studies **classes** of languages and their computational properties
 - ▶ Regular language (this lecture)
 - ▶ Context free language (next lecture)
- Main goal is to solve the **membership problem**
 - ▶ Whether a string is in a language or not
- How? By defining its **grammar**

Examples

- Binary strings that start with 0 and end with 1
 - ▶ $\{ 01, 001, 011, 0001, \dots \}$ ✓
 - ▶ $\{ 1, 0, 00, 11, 100, \dots \}$ ✗
- Even-length sequences from alphabet $\{a, b\}$
 - ▶ $\{ aa, ab, ba, bb, aaaa, \dots \}$ ✓
 - ▶ $\{ aaa, aba, bbb, \dots \}$ ✗
- English sentences that start with *wh*-word and end in ?
 - ▶ $\{ \textit{what} ?, \textit{where my pants} ?, \dots \}$ ✓

Beyond Membership Problem...

- Membership
 - ▶ Is the string part of the language? Y/N
- Scoring
 - ▶ Graded membership
 - ▶ “How acceptable is a string?” (language models!)
- Transduction
 - ▶ “Translate” one string into another (stemming!)

Overview

- Regular languages
 - Finite state acceptors & transducers
 - Modelling word morphology
- allow for
com
trans*
- allow for
com
trans*
- allow for
com
definitions.*

Regular Languages

- Regular language: the simplest class of languages
- Any **regular expression** is a regular language
 - ▶ Describes what strings are part of the language

Regular Languages

- Formally, a regular expression includes the following operations:
 - ▶ Symbol drawn from alphabet, Σ *all symbols*
 - ▶ Empty string, ϵ
 - ▶ Concatenation of two regular expressions, RS
 - ▶ Alternation of two regular expressions, $R|S$
 - ▶ Kleene star for 0 or more repeats, R^*
 - ▶ Parenthesis () to define scope of operations

Examples of Regular Languages

- Binary strings that start with 0 and end with 1
 - ▶ $0(0|1)^*1$
- Even-length sequences from alphabet $\{a, b\}$
 - ▶ $((aa)|(ab)|(ba)|(bb))^*$
- English sentences that start with *wh*-word and end in ?
 - ▶ $((what)|(where)|(why)|(which)|(whose)|(whom)) \Sigma^* ?$

. *

Properties of Regular Languages

- Closure: if we take regular languages L_1 and L_2 and merge them, is the resulting language regular?
Closed Operation
- RLs are closed under the following:
 - ▶ concatenation and union — follows from definition
 - ▶ intersection: strings that are valid in both L_1 and L_2
 - ▶ negation: strings that are not in L
- Extremely versatile! Can have RLs for different properties of language, and use them together
 - ▶ core algorithms will still apply

Finite State Acceptors

- Regular expression defines a regular language
- But it doesn't give an algorithm to check whether a string belongs to the language
- **Finite state acceptors** (FSA) describes the computation involved for membership checking

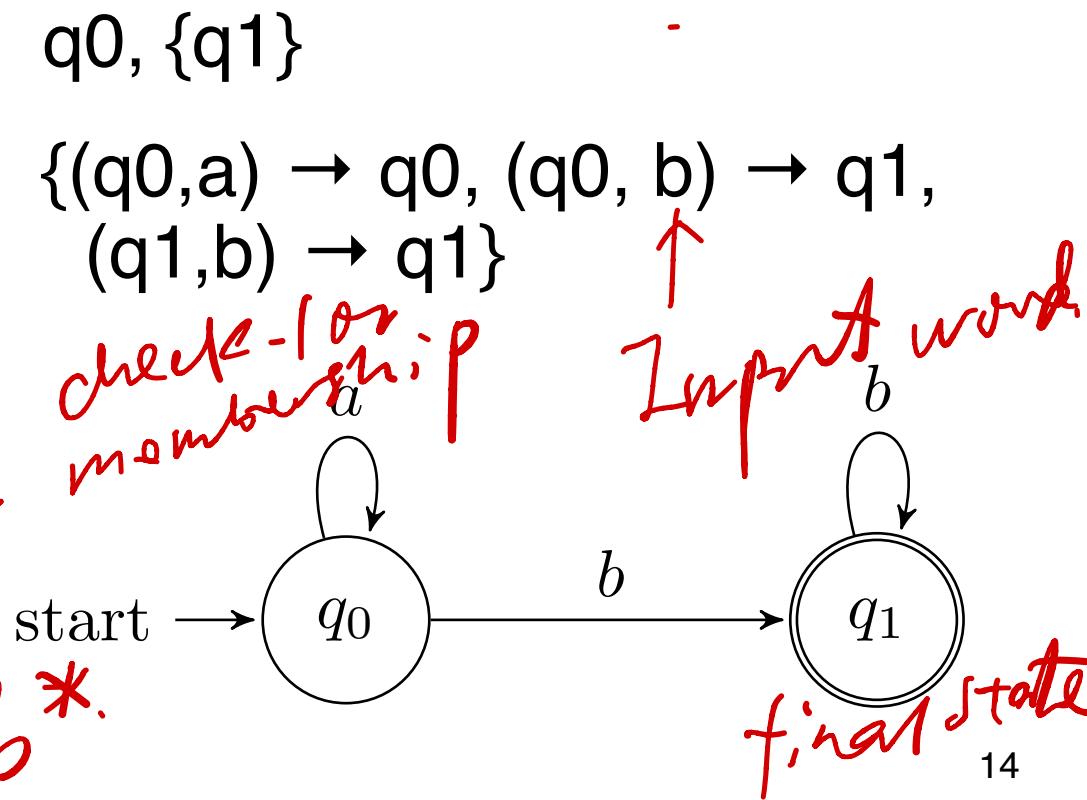
Finite State Acceptors

- FSA consists:
 - ▶ alphabet of input symbols, Σ
 - ▶ set of states, Q
 - ▶ start state, $q_0 \in Q$
 - ▶ final states, $F \subseteq Q$
 - ▶ transition function
 - symbol and state \rightarrow next state
 - Accepts strings if there is **path** from q_0 to a final state with ^{valid-} transitions matching each symbol
- Djisktra's shortest-path algorithm, $O(V \log V + E)$

Example FSA

- Input alphabet $\{a, b\}$
- States $\{q_0, q_1\}$
- Start, final states $q_0, \{q_1\}$
- Transition function $\{(q_0, a) \rightarrow q_0, (q_0, b) \rightarrow q_1, (q_1, b) \rightarrow q_1\}$
- Note: seeing a in q_1 results in failure
- Accepts a^*bb^*

a^*b^* $a^* \rightarrow b \rightarrow b^*$



Derivational Morphology

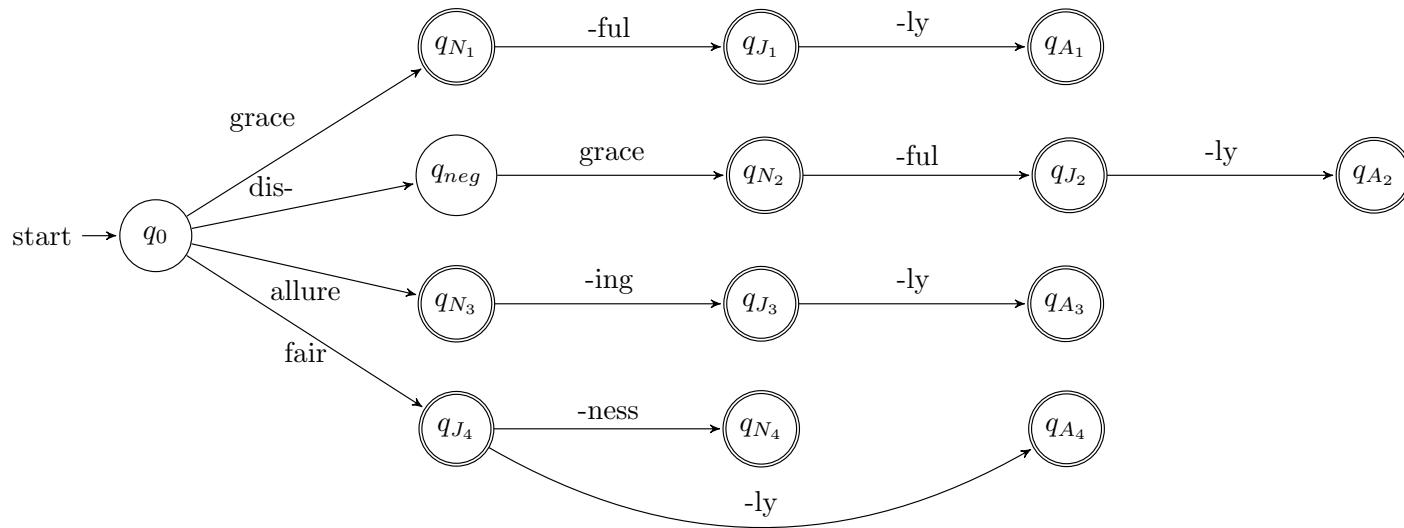
- Use of affixes to change word to another grammatical category
- *grace* → *graceful* → *gracefully*
- *grace* → *disgrace* → *disgracefully*
- *allure* → *alluring* → *alluringly*
- *allure* → **allureful*
- *allure* → **disallure*

FSA for Morphology

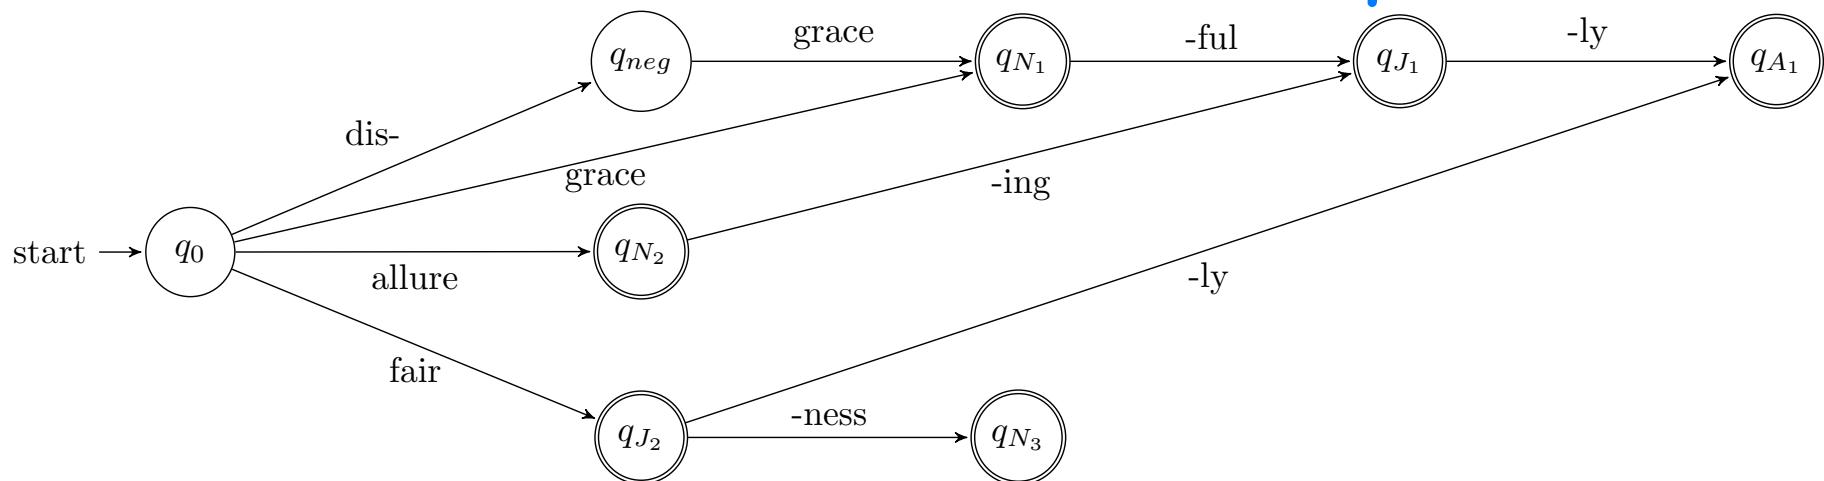
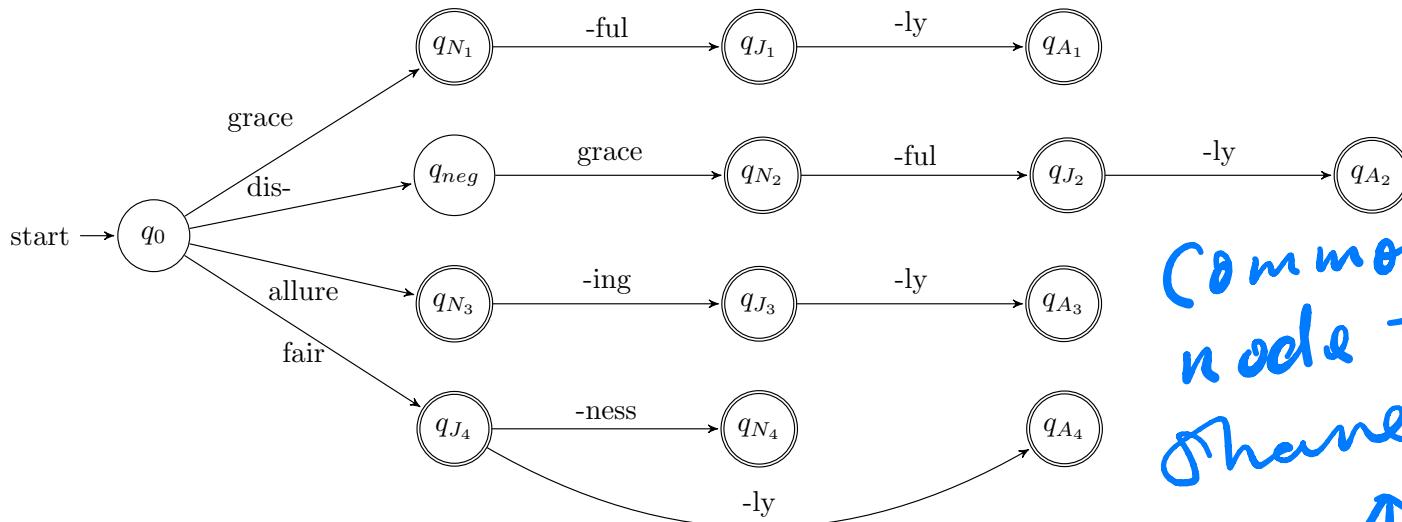
- (Fairly) consistent process— can we describe this as a regular language?
 - ▶ want to accept valid forms, and reject invalid ones [flagged with *]
 - ▶ generalise to other words, e.g., nouns that behave like *grace* or *allure*

Reg L \rightarrow P Morphology.

FSA for Word Morphology



FSA for Word Morphology



Weighted FSA

Weighted FSA

- Some words are more **possible** than others
 - ▶ *fishful* vs. *disgracelyful*
 - ▶ *musicky* vs. *writey*
- Graded measure of acceptability — weighted FSA adds/changes the following:

- ▶ set of states Q
 - ▶ alphabet of input symbols Σ
 - ▶ start state weight function, $\lambda: Q \rightarrow \mathbb{R}$
 - ▶ final state weight function, $\rho: Q \rightarrow \mathbb{R}$
 - ▶ transition function, $\delta: (Q, \Sigma, Q) \rightarrow \mathbb{R}$
- weights for
start/ing state &
transition func'.*
- input output weight.*

WFSA Shortest-Path

- Total score of a path $\pi = t_1, \dots, t_N$ now

$$\lambda(t_0) + \sum_{i=1}^N \delta(t_i) + \rho(t_N)$$

→ all paths took
↓ final

each t is an edge, so more formally using from &/or to states and edge label in score calculation

- Use *shortest-path algorithm* to find π with min. cost
 - $O(V \log V + E)$, as before

Dijkstra

N-gram LMs as WFSA

- Recall LM calculates score of string as follows

$$p(w_1, \dots, w_M) \approx \prod_{m=1}^M p_n(w_m | w_{m-1}, \dots, w_{m-n+1}).$$

- Unigram language model:

- One state: q_0
- State transition score: $\delta(q_0, \omega, q_0) = \log p_1(\omega)$
- Initial and final state scores = 0.
- Path score for w_1, w_2, \dots, w_M :

$$0 + \sum_m^M \delta(q_0, w_m, q_0) + 0 = \sum_m^M \log p_1(w_m).$$

Bigram LM

- Bigram sentence probability:

$$P(w_1, w_2, \dots, w_M) = \prod_{i=1}^M P(w_i | w_{i-1}) \quad (\text{bigram})$$

- Implemented as WFSA

► Σ = set of word types

► $Q = \Sigma$ (no. of states = no. of word types)

$$\delta(q_i, \omega, q_j) = \begin{cases} \log \Pr(w_m = j | w_{m-1} = i), & \omega = j \\ -\infty, & \omega \neq j \end{cases}$$

$$\lambda(q_i) = \log \Pr(w_1 = i | w_0 = \square)$$

$$\rho(q_i) = \log \Pr(w_{M+1} = \blacksquare | w_M = i).$$

start symbol.
end symbol.

need one state for each word.
need to remember the previous one.

1 state to remember
the previous one.

do not allow
model to
transition to
the
new state

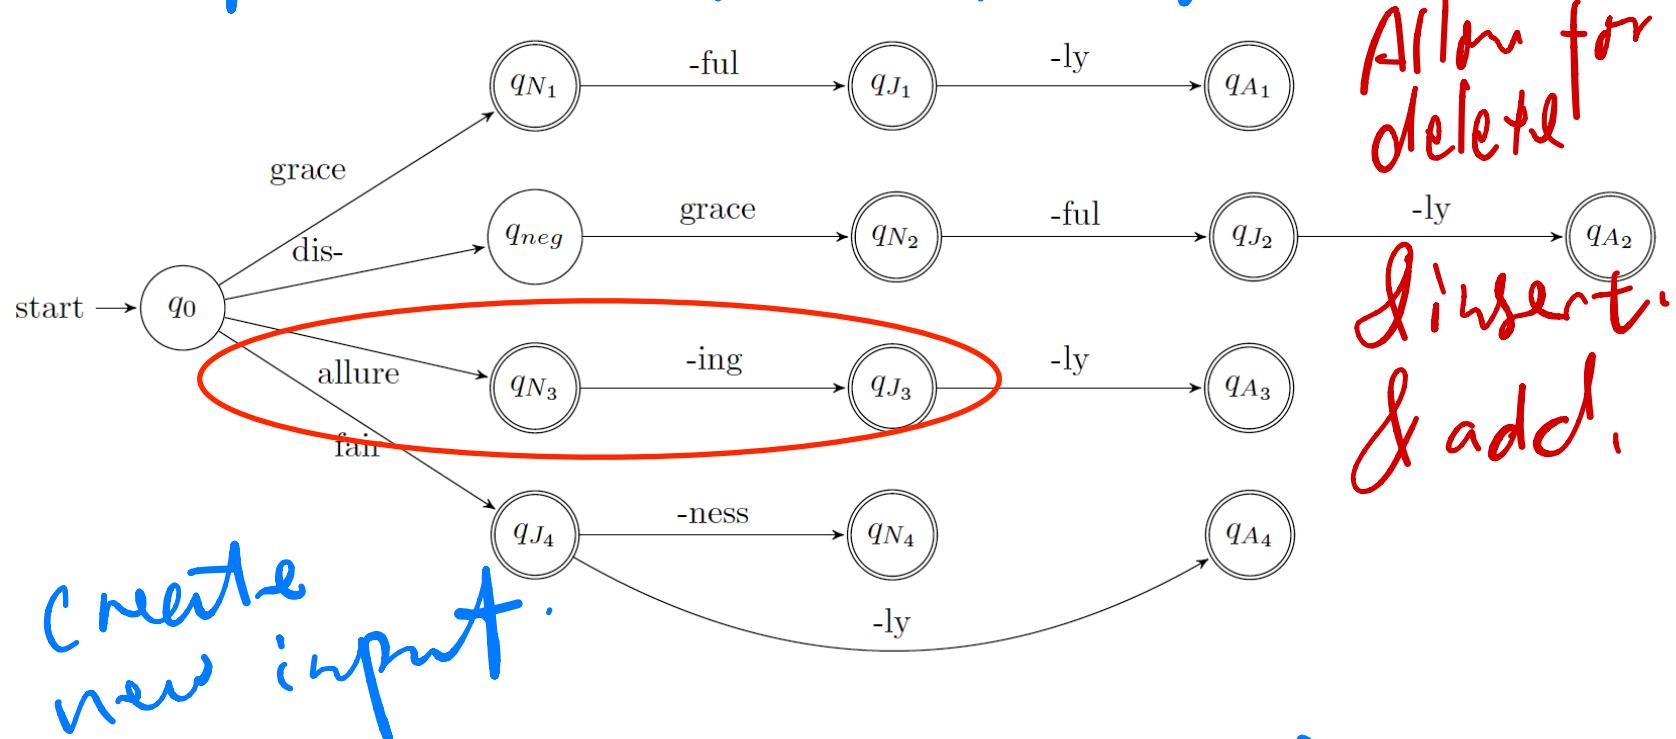
Finite State Transducer

Finite State Transducers (FST)

- Often don't want to just accept or score strings
 - ▶ want to translate them into another language, correct grammar, parse their structure, etc

FSA for Word Morphology

Derivations of morphology.



FSA: allure + ing = alluring X.

FST: allure + ing = alluring

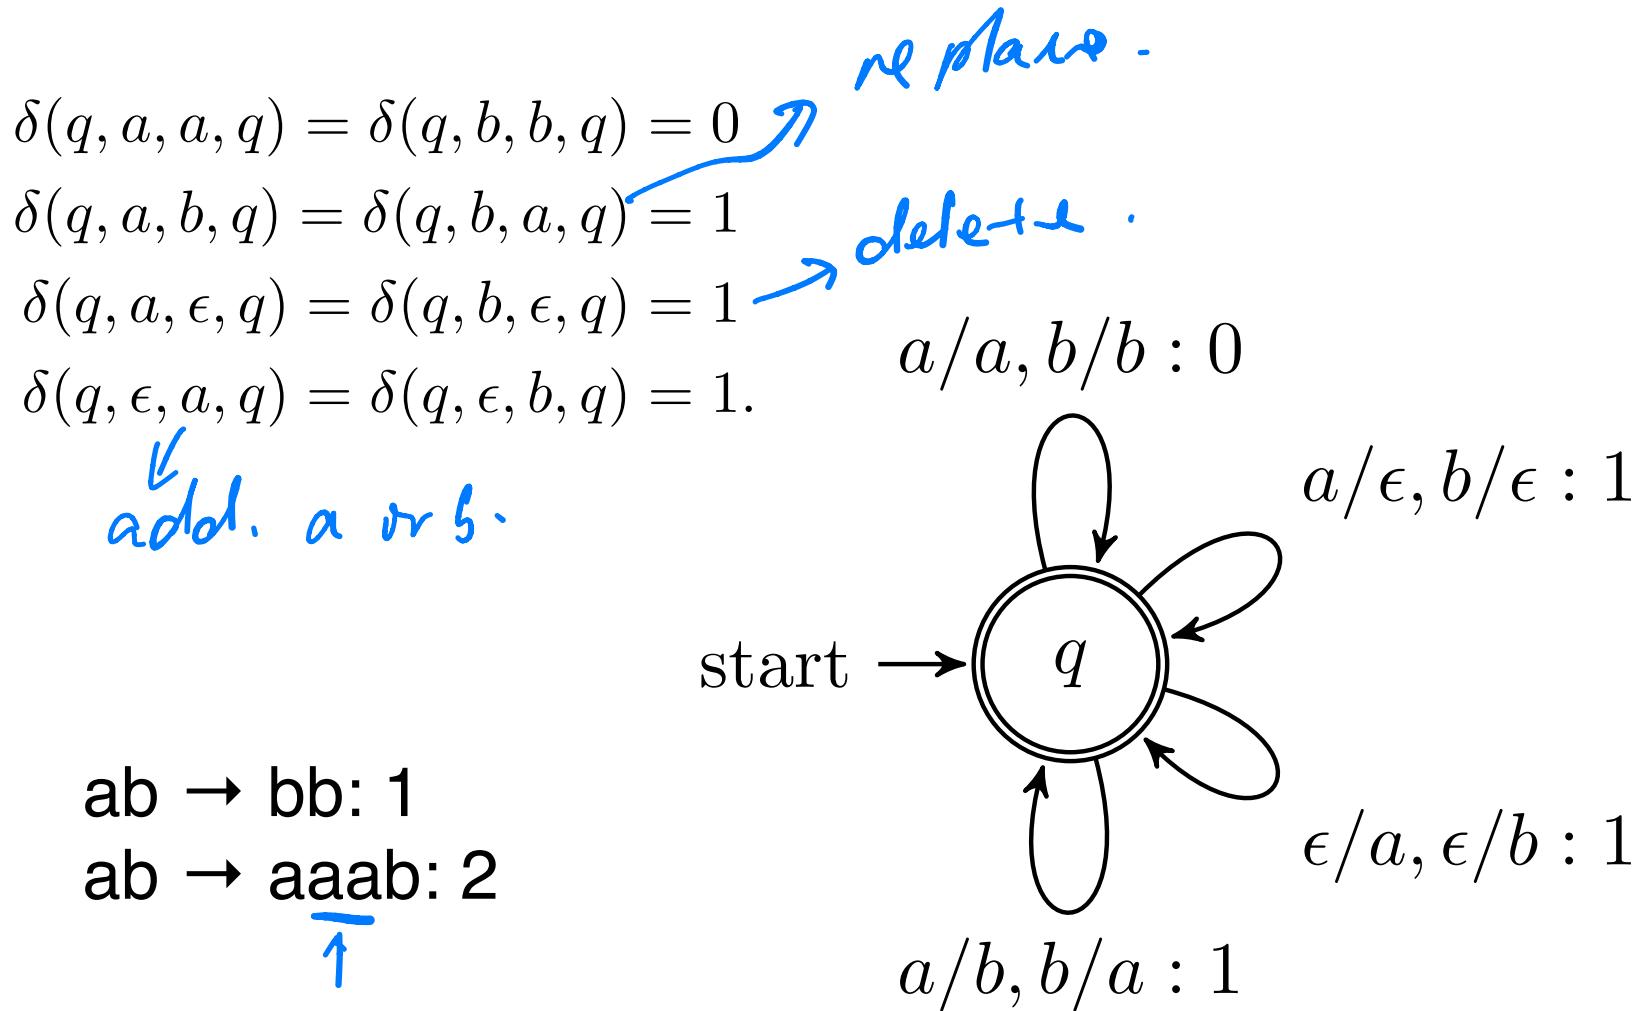
Transducer.

Finite State Transducers

input alphabet

- FST add string output capability to FSAs
 - includes an *output alphabet*
 - and transitions now take input symbol and **emit output symbol** (Q, Σ, Σ, Q)
- Can be **weighted** = WFST
 - Graded scores for transition
- E.g., edit distance as **WFST** which takes one string, and outputs the other
 - zero cost only if strings are identical

Edit Distance Automata



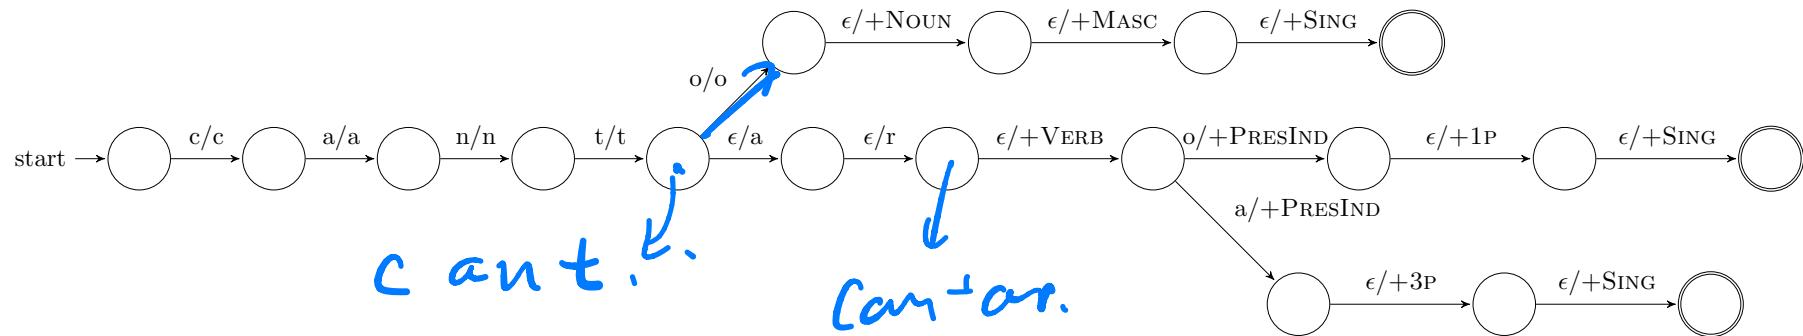
FST for Inflectional Morphology

- Verb inflections in Spanish must match the subject in person & number
- Goal of **morphological analysis**:
 - canto* → *cantar+VERB+present+1P+singular*

	cantar	to sing
1P singular	yo canto	I sing
2P singular	tu cantas	you sing
3P singular	ella canta	she sings
1P plural	nostotros cantamos	we sing
2P plural	vosotros cantáis	you sing
3P plural	ellas cantan	they sing

Inf
Morph

FST for Spanish Inflection



canto → *canto+Noun+Masc+Sing*

canto → *cantar+Verb+PresInd+1P+Sing*

FST Composition

- Compose two FSTs by taking output of one FST, T₁, and giving this as input to FST T₂
 - ▶ denoted $T_1 \circ T_2$; and results in another FST
 - ▶ can also compose FST with FSA, resulting in a FST

$T_1 \circ T_2$

$FST_1 \rightarrow T_1 \rightarrow FST_2 \rightarrow T_2.$

FST Composition Example

- Allows development of different processes as FSTs:
 - ▶ -ed added to signal past tense in English
 - ▶ *cook* → *cooked*, *want* → *wanted*
 - ▶ But when the word ends with 'e', then we want to avoid producing a spelling with consecutive e's (*bakeed*)
 - ▶ Solution: build 2 FSTs
 - ▶ FST T1: *bake+PAST* → *bake+ed*
 - ▶ FST T2: *bake+ed* → *baked*

wednes /ə/

Is Natural Language Regular?

Sometimes...

- Example:

the mouse that ran.

the cat that killed the mouse that ran.

...

the lion that bullied the hyena that bit the dog that chased the cat that killed the mouse that ran

...

- Length is unbounded (*recursive*), but structure is local → can describe with FSA = Regular
- $(\text{Det Noun Prep Verb})^*$

Non-Regular Languages

FSA.

- Arithmetic expressions with balanced parentheses

▶ $(a + (b \times (c / d)))$

▶ Can have ~~arbitrarily many opening parentheses~~

▶ Need to remember how many open parentheses, to produce the same number of closed parentheses

▶ Can't be done with ~~finite number of states~~.

not be
able
to
do.

- $a^n b^n \rightarrow$ not a regular language.

cannot be captured $\Rightarrow 'n'$

Center Embedding

- Center embedding of relative clauses
 - ▶ The cat loves Mozart
 - ▶ The cat **the dog chased** loves Mozart
 - ▶ The cat **the dog the rat bit chased** loves Mozart
 - ▶ The cat **the dog the rat the elephant admired bit chased** loves Mozart
 - Need to remember the n subject nouns, to ensure n verbs follow (and that they agree etc)
 - Requires (at least) **context-free grammar** (next lecture!)
-
- aⁿ bⁿ
- n subject nouns
- not regular

Summary

- Concept of a language, and grammar
- Regular languages
- Finite state automata: acceptors, transducers
- Closure properties
- Weighted variants & shortest path inference
- Application to edit distance, morphology

Reading

- Reading
 - ▶ E18, Chapter 9.1